

DSC 475 - Project 2

Hao Ma

1.1 (a), (b)

Note: The two plots are embedded in accuracy&loss reports.

```
In [1]: import torch
import torch.nn as nn
import torch.nn.functional as F
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```

In [14]: class Net(nn.Module):
    # network structure for 1.1
    def __init__(self):
        super(Net, self).__init__()
        # three hidden layers with 20 neurons each
        self.fc1 = nn.Linear(2, 20)
        self.fc2 = nn.Linear(20, 20)
        self.fc3 = nn.Linear(20, 20)
        self.fc4 = nn.Linear(20, 20)
        self.fc5 = nn.Linear(20, 2)

    def forward(self, x):
        # choice of activation function: relu
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = F.relu(self.fc3(x))
        x = F.relu(self.fc4(x))
        x = self.fc5(x)
        return F.log_softmax(x)

def plot_decision_boundary(net, X, y):
    # A function to plot the decision boundary
    x_min, x_max = -0.5, 1.5
    y_min, y_max = -0.5, 1.5
    h = 0.01
    # Generate a grid of points with distance h between them
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

    # Predict the function value for the whole grid
    X_out = net(torch.tensor(np.c_[xx.ravel(), yy.ravel()], dtype = torch.float))

    Z = X_out.data.max(1)[1]
    Z = Z.reshape(xx.shape)
    # Plot the contour and training examples
    plt.contourf(xx, yy, Z, cmap=plt.cm.Spectral)
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Spectral, s = 1)

# XOR data (columns names have been added in the csv file)
data = pd.read_csv("XOR.csv")
X = data.values[:, 0:2] # Take only the first two features.
X = torch.tensor(X, dtype = torch.float)
y = data.values[:, 2]
y = torch.tensor(y, dtype = torch.long)

# Training
net = Net()
# create a stochastic gradient descent optimizer
learning_rate = 0.05 # for faster convergence
optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate, momentum=0.9)

# create a loss function
criterion = nn.NLLLoss()

nepochs = 10000
data, target = X, y

```

```

earliest = True
# run the main training loop
for epoch in range(nepochs):
    optimizer.zero_grad()
    # forward propagate
    net_out = net(data)
    # compute loss
    loss = criterion(net_out, target)
    # backpropagate
    loss.backward()
    # update parameters
    optimizer.step()
    # print out report until 100% accuracy is achieved or
    # the loss falls below 0.0001
    print('Epoch ', epoch, 'Loss ', loss.item())
    pred = net_out.data.max(1)[1] # get the index of the max log-probability
    correctidx = pred.eq(target.data)
    ncorrect = correctidx.sum()
    accuracy = ncorrect.item()/len(data)
    print('Training accuracy is ', accuracy)

    if accuracy == 1 and earliest:
        # As required, plot the decision boundaries of the earliest network in
        the training process
        # that achieves 100% accuracy
        earliest = False
        print('**100% accuracy is achieved here**')
        plot_decision_boundary(net, X, y)
        plt.show()

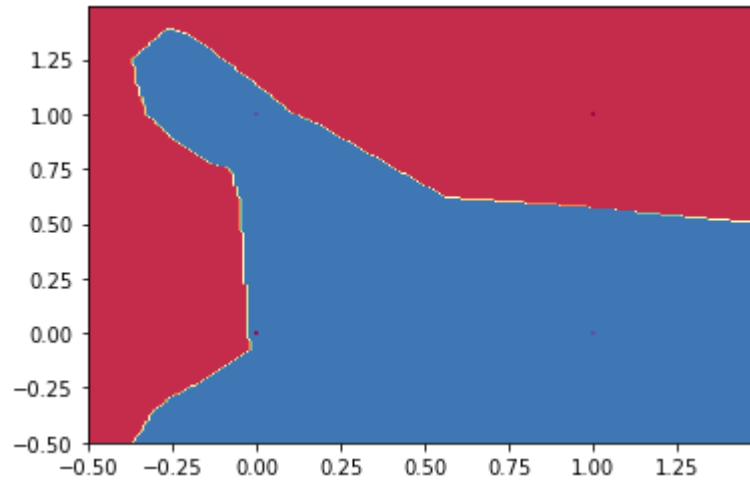
    if loss.item() < 0.0001:
        # As required, plot the decision boundaries after the loss falls below
        0.0001
        print('**Loss falls below 0.0001 here**')
        plot_decision_boundary(net, X, y)
        plt.show()
        break

```

```
C:\Anaconda3\lib\site-packages\ipykernel_launcher.py:19: UserWarning: Implicit dimension choice for log_softmax has been deprecated. Change the call to include dim=X as an argument.
```

Epoch 0 Loss 0.6945462226867676
Training accuracy is 0.5
Epoch 1 Loss 0.694356381893158
Training accuracy is 0.5
Epoch 2 Loss 0.6940532326698303
Training accuracy is 0.5
Epoch 3 Loss 0.6937108039855957
Training accuracy is 0.5
Epoch 4 Loss 0.6933308839797974
Training accuracy is 0.5
Epoch 5 Loss 0.6929678916931152
Training accuracy is 0.5
Epoch 6 Loss 0.6926510334014893
Training accuracy is 0.5
Epoch 7 Loss 0.6923943758010864
Training accuracy is 0.5
Epoch 8 Loss 0.692197859287262
Training accuracy is 0.75
Epoch 9 Loss 0.6920502185821533
Training accuracy is 0.5
Epoch 10 Loss 0.6919325590133667
Training accuracy is 0.5
Epoch 11 Loss 0.6918235421180725
Training accuracy is 0.5
Epoch 12 Loss 0.6917003393173218
Training accuracy is 0.5
Epoch 13 Loss 0.6915504336357117
Training accuracy is 0.5
Epoch 14 Loss 0.6913648247718811
Training accuracy is 0.5
Epoch 15 Loss 0.6911377310752869
Training accuracy is 0.5
Epoch 16 Loss 0.6909124255180359
Training accuracy is 0.5
Epoch 17 Loss 0.6906532049179077
Training accuracy is 0.5
Epoch 18 Loss 0.6903915405273438
Training accuracy is 0.5
Epoch 19 Loss 0.6900997161865234
Training accuracy is 0.5
Epoch 20 Loss 0.6897867918014526
Training accuracy is 0.5
Epoch 21 Loss 0.6894621849060059
Training accuracy is 0.5
Epoch 22 Loss 0.6891050934791565
Training accuracy is 0.75
Epoch 23 Loss 0.6887248158454895
Training accuracy is 0.75
Epoch 24 Loss 0.6883214116096497
Training accuracy is 0.75
Epoch 25 Loss 0.6880104541778564
Training accuracy is 0.75
Epoch 26 Loss 0.6876884698867798
Training accuracy is 0.75
Epoch 27 Loss 0.6873420476913452
Training accuracy is 0.75
Epoch 28 Loss 0.6869895458221436


Training accuracy is 0.75
Epoch 29 Loss 0.6866284012794495
Training accuracy is 0.75
Epoch 30 Loss 0.6862086057662964
Training accuracy is 0.75
Epoch 31 Loss 0.6857434511184692
Training accuracy is 0.75
Epoch 32 Loss 0.6852546334266663
Training accuracy is 0.75
Epoch 33 Loss 0.6847296953201294
Training accuracy is 0.75
Epoch 34 Loss 0.6841648817062378
Training accuracy is 0.75
Epoch 35 Loss 0.6835817098617554
Training accuracy is 0.75
Epoch 36 Loss 0.6829493641853333
Training accuracy is 0.75
Epoch 37 Loss 0.6822612285614014
Training accuracy is 1.0
100% accuracy is achieved here



Epoch 38 Loss 0.6815095543861389
Training accuracy is 0.75
Epoch 39 Loss 0.6806870698928833
Training accuracy is 0.75
Epoch 40 Loss 0.6797855496406555
Training accuracy is 0.75
Epoch 41 Loss 0.6789296865463257
Training accuracy is 0.75
Epoch 42 Loss 0.6778762340545654
Training accuracy is 0.75
Epoch 43 Loss 0.676721453666687
Training accuracy is 0.75
Epoch 44 Loss 0.6755290627479553
Training accuracy is 0.75
Epoch 45 Loss 0.6741771697998047
Training accuracy is 0.75
Epoch 46 Loss 0.6726700663566589
Training accuracy is 0.75
Epoch 47 Loss 0.6710194945335388
Training accuracy is 0.75
Epoch 48 Loss 0.6692386269569397
Training accuracy is 0.75
Epoch 49 Loss 0.6672821044921875
Training accuracy is 0.75
Epoch 50 Loss 0.6650102138519287
Training accuracy is 0.75
Epoch 51 Loss 0.6625635027885437
Training accuracy is 0.75
Epoch 52 Loss 0.6598827242851257
Training accuracy is 0.75
Epoch 53 Loss 0.6569311618804932
Training accuracy is 0.75
Epoch 54 Loss 0.6536464095115662
Training accuracy is 0.75
Epoch 55 Loss 0.6499893665313721
Training accuracy is 0.75
Epoch 56 Loss 0.6459165811538696
Training accuracy is 0.75
Epoch 57 Loss 0.6410896182060242
Training accuracy is 0.75
Epoch 58 Loss 0.6355979442596436
Training accuracy is 0.75
Epoch 59 Loss 0.6294361352920532
Training accuracy is 0.75
Epoch 60 Loss 0.6226579546928406
Training accuracy is 0.75
Epoch 61 Loss 0.616114616394043
Training accuracy is 0.75
Epoch 62 Loss 0.607833206653595
Training accuracy is 0.75
Epoch 63 Loss 0.5985419750213623
Training accuracy is 0.75
Epoch 64 Loss 0.5876613259315491
Training accuracy is 0.75
Epoch 65 Loss 0.5751106142997742
Training accuracy is 0.75
Epoch 66 Loss 0.560251772403717

Training accuracy is 0.75
Epoch 67 Loss 0.5439274907112122
Training accuracy is 0.75
Epoch 68 Loss 0.5251119136810303
Training accuracy is 0.75
Epoch 69 Loss 0.5049505233764648
Training accuracy is 0.75
Epoch 70 Loss 0.4819396436214447
Training accuracy is 0.75
Epoch 71 Loss 0.456208199262619
Training accuracy is 1.0
Epoch 72 Loss 0.42637455463409424
Training accuracy is 1.0
Epoch 73 Loss 0.393406480550766
Training accuracy is 1.0
Epoch 74 Loss 0.35736820101737976
Training accuracy is 1.0
Epoch 75 Loss 0.31923383474349976
Training accuracy is 1.0
Epoch 76 Loss 0.27747732400894165
Training accuracy is 1.0
Epoch 77 Loss 0.23408086597919464
Training accuracy is 1.0
Epoch 78 Loss 0.18903197348117828
Training accuracy is 1.0
Epoch 79 Loss 0.14502406120300293
Training accuracy is 1.0
Epoch 80 Loss 0.10646328330039978
Training accuracy is 1.0
Epoch 81 Loss 0.07477398216724396
Training accuracy is 1.0
Epoch 82 Loss 0.05193035304546356
Training accuracy is 1.0
Epoch 83 Loss 0.035627152770757675
Training accuracy is 1.0
Epoch 84 Loss 0.024248354136943817
Training accuracy is 1.0
Epoch 85 Loss 0.01676982454955578
Training accuracy is 1.0
Epoch 86 Loss 0.011914549395442009
Training accuracy is 1.0
Epoch 87 Loss 0.008497071452438831
Training accuracy is 1.0
Epoch 88 Loss 0.006136741489171982
Training accuracy is 1.0
Epoch 89 Loss 0.004476932343095541
Training accuracy is 1.0
Epoch 90 Loss 0.0033141691237688065
Training accuracy is 1.0
Epoch 91 Loss 0.002492073690518737
Training accuracy is 1.0
Epoch 92 Loss 0.0019027431262657046
Training accuracy is 1.0
Epoch 93 Loss 0.0014740444021299481
Training accuracy is 1.0
Epoch 94 Loss 0.001162998960353434
Training accuracy is 1.0

Epoch 95 Loss 0.0009316123323515058
Training accuracy is 1.0
Epoch 96 Loss 0.0007578188669867814
Training accuracy is 1.0
Epoch 97 Loss 0.0006256877677515149
Training accuracy is 1.0
Epoch 98 Loss 0.0005241867038421333
Training accuracy is 1.0
Epoch 99 Loss 0.00044528322177939117
Training accuracy is 1.0
Epoch 100 Loss 0.0003832040820270777
Training accuracy is 1.0
Epoch 101 Loss 0.0003339606919325888
Training accuracy is 1.0
Epoch 102 Loss 0.00029463559621945024
Training accuracy is 1.0
Epoch 103 Loss 0.0002626081695780158
Training accuracy is 1.0
Epoch 104 Loss 0.00023644912289455533
Training accuracy is 1.0
Epoch 105 Loss 0.00021484798344317824
Training accuracy is 1.0
Epoch 106 Loss 0.00019685158622451127
Training accuracy is 1.0
Epoch 107 Loss 0.0001817450684029609
Training accuracy is 1.0
Epoch 108 Loss 0.0001689326309133321
Training accuracy is 1.0
Epoch 109 Loss 0.00015808662283234298
Training accuracy is 1.0
Epoch 110 Loss 0.00014876014029141515
Training accuracy is 1.0
Epoch 111 Loss 0.0001407446397934109
Training accuracy is 1.0
Epoch 112 Loss 0.00013380181917455047
Training accuracy is 1.0
Epoch 113 Loss 0.0001277528645005077
Training accuracy is 1.0
Epoch 114 Loss 0.00012241902004461735
Training accuracy is 1.0
Epoch 115 Loss 0.00011780030035879463
Training accuracy is 1.0
Epoch 116 Loss 0.00011374773748684675
Training accuracy is 1.0
Epoch 117 Loss 0.00011017193173756823
Training accuracy is 1.0
Epoch 118 Loss 0.00010695368837332353
Training accuracy is 1.0
Epoch 119 Loss 0.00010409301467007026
Training accuracy is 1.0
Epoch 120 Loss 0.00010150053276447579
Training accuracy is 1.0
Epoch 121 Loss 9.920603042701259e-05
Training accuracy is 1.0
Loss falls below 0.0001 here



1.2 (a), (b)

```

In [16]: class Net(nn.Module):
    # network structure for 1.2: 1 hidden layer with 5 neurons
    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(2, 5)
        self.fc2 = nn.Linear(5, 5)
        self.fc3 = nn.Linear(5, 2)

    def forward(self, x):
        # choice of activation function: relu
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return F.log_softmax(x)

net = Net()
learning_rate = 0.05 # for faster convergence
optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate, momentum=0.9)
criterion = nn.NLLLoss()

nepochs = 30000
data, target = X, y
earliest = True
for epoch in range(nepochs):
    optimizer.zero_grad()
    # forward propagate
    net_out = net(data)
    # compute loss
    loss = criterion(net_out, target)
    # backpropagate
    loss.backward()
    # update parameters
    optimizer.step()

    pred = net_out.data.max(1)[1]
    correctidx = pred.eq(target.data)
    ncorrect = correctidx.sum()
    accuracy = ncorrect.item()/len(data)

    if accuracy == 1 and earliest:
        # As required, plot the decision boundaries of the earliest network in
        the training process
        # that achieves 100% accuracy
        earliest = False
        print('**100% accuracy is achieved here**')
        print('Epoch ', epoch, 'Loss ', loss.item())
        print('Training accuracy is ', accuracy)
        plot_decision_boundary(net, X, y)
        plt.show()

    if loss.item() < 0.0001:
        # As required, plot the decision boundaries after the loss falls below
        0.0001
        print('**Loss falls below 0.0001 here**')
        print('Epoch ', epoch, 'Loss ', loss.item())

```


```
print('Training accuracy is ', accuracy)
plot_decision_boundary(net, X, y)
plt.show()
break
```

****100% accuracy is achieved here****

Epoch 24 Loss 0.6558001041412354

Training accuracy is 1.0


C:\Anaconda3\lib\site-packages\ipykernel_launcher.py:14: UserWarning: Implicit dimension choice for log_softmax has been deprecated. Change the call to include dim=X as an argument.



****Loss falls below 0.0001 here****

Epoch 1730 Loss 9.994937136070803e-05

Training accuracy is 1.0




The smallest network that can achieve 100% accuracy (with reasonable number of iterations) contains 1 hidden layer with 5 neurons.

2.1

```
In [22]: mydata = pd.read_csv('FeedForward_Data_ellipse.csv', header=None)
X = mydata.values[:, 0:2] # features
X = torch.tensor(X, dtype = torch.float)
y = mydata.values[:, 2] # labels
y = torch.tensor(y, dtype = torch.long)

# plot the original data
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Spectral, s = 1)
```

Out[22]: <matplotlib.collections.PathCollection at 0x29df2874508>



```

In [24]: class Net(nn.Module):
    # network structure for 2.1
    def __init__(self):
        super(Net, self).__init__()
        # 10 hidden layers with 20 neurons each
        self.fc1 = nn.Linear(2, 20)
        self.fc2 = nn.Linear(20, 20)
        self.fc3 = nn.Linear(20, 20)
        self.fc4 = nn.Linear(20, 20)
        self.fc5 = nn.Linear(20, 20)
        self.fc6 = nn.Linear(20, 20)
        self.fc7 = nn.Linear(20, 20)
        self.fc8 = nn.Linear(20, 20)
        self.fc9 = nn.Linear(20, 20)
        self.fc10 = nn.Linear(20, 20)
        self.fc11 = nn.Linear(20, 20)
        self.fc12 = nn.Linear(20, 2)

    def forward(self, x):
        # choice of activation function: relu
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = F.relu(self.fc3(x))
        x = F.relu(self.fc4(x))
        x = F.relu(self.fc5(x))
        x = F.relu(self.fc6(x))
        x = F.relu(self.fc7(x))
        x = F.relu(self.fc8(x))
        x = F.relu(self.fc9(x))
        x = F.relu(self.fc10(x))
        x = F.relu(self.fc11(x))
        x = self.fc12(x)
        return F.log_softmax(x)

def plot_decision_boundary(net, X, y):
    # A function to plot the decision boundary
    x_min, x_max = -1.0, 1.0
    y_min, y_max = -1.0, 1.0
    h = 0.01
    # Generate a grid of points with distance h between them
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
    # Predict the function value for the whole grid
    X_out = net(torch.tensor(np.c_[xx.ravel(), yy.ravel()], dtype = torch.float))
    Z = X_out.data.max(1)[1]
    Z = Z.reshape(xx.shape)
    # Plot the contour and training examples
    plt.contourf(xx, yy, Z, cmap=plt.cm.Spectral)
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Spectral, s = 1)

# Training
net = Net()
learning_rate = 0.1

```

```

optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate, momentum=0.9)
criterion = nn.NLLLoss()

nepochs = 5001
data, target = X, y
for epoch in range(nepochs):
    optimizer.zero_grad()
    # forward propagate
    net_out = net(data)
    # compute loss
    loss = criterion(net_out, target)
    # backpropagate
    loss.backward()
    # update parameters
    optimizer.step()

    pred = net_out.data.max(1)[1]
    correctidx = pred.eq(target.data)
    ncorrect = correctidx.sum()
    accuracy = ncorrect.item()/len(data)

    if epoch % 50 == 0:
        # print out report every 50 iterations
        print('Epoch ', epoch, 'Loss ', loss.item())
        print('Training accuracy is ', accuracy)

```

```
C:\Anaconda3\lib\site-packages\ipykernel_launcher.py:34: UserWarning: Implicit dimension choice for log_softmax has been deprecated. Change the call to include dim=X as an argument.
```


Epoch 0 Loss 0.7107270956039429
Training accuracy is 0.4434344951923077
Epoch 50 Loss 0.6867424845695496
Training accuracy is 0.5565655048076923
Epoch 100 Loss 0.6867023706436157
Training accuracy is 0.5565655048076923
Epoch 150 Loss 0.686763346195221
Training accuracy is 0.5565655048076923
Epoch 200 Loss 0.6866981983184814
Training accuracy is 0.5565655048076923
Epoch 250 Loss 0.686698853969574
Training accuracy is 0.5565655048076923
Epoch 300 Loss 0.6866989135742188
Training accuracy is 0.5565655048076923
Epoch 350 Loss 0.686698853969574
Training accuracy is 0.5565655048076923
Epoch 400 Loss 0.686698853969574
Training accuracy is 0.5565655048076923
Epoch 450 Loss 0.686698853969574
Training accuracy is 0.5565655048076923
Epoch 500 Loss 0.6866987943649292
Training accuracy is 0.5565655048076923
Epoch 550 Loss 0.6866986751556396
Training accuracy is 0.5565655048076923
Epoch 600 Loss 0.6866986751556396
Training accuracy is 0.5565655048076923
Epoch 650 Loss 0.6866985559463501
Training accuracy is 0.5565655048076923
Epoch 700 Loss 0.6866984963417053
Training accuracy is 0.5565655048076923
Epoch 750 Loss 0.6866984963417053
Training accuracy is 0.5565655048076923
Epoch 800 Loss 0.6866983771324158
Training accuracy is 0.5565655048076923
Epoch 850 Loss 0.6866983771324158
Training accuracy is 0.5565655048076923
Epoch 900 Loss 0.686698317527771
Training accuracy is 0.5565655048076923
Epoch 950 Loss 0.686698317527771
Training accuracy is 0.5565655048076923
Epoch 1000 Loss 0.686698317527771
Training accuracy is 0.5565655048076923
Epoch 1050 Loss 0.6866982579231262
Training accuracy is 0.5565655048076923
Epoch 1100 Loss 0.6866981983184814
Training accuracy is 0.5565655048076923
Epoch 1150 Loss 0.6866981983184814
Training accuracy is 0.5565655048076923
Epoch 1200 Loss 0.6866980791091919
Training accuracy is 0.5565655048076923
Epoch 1250 Loss 0.6866980791091919
Training accuracy is 0.5565655048076923
Epoch 1300 Loss 0.6866980791091919
Training accuracy is 0.5565655048076923
Epoch 1350 Loss 0.6866980791091919
Training accuracy is 0.5565655048076923
Epoch 1400 Loss 0.6866980791091919

Training accuracy is 0.5565655048076923
Epoch 1450 Loss 0.6866980195045471
Training accuracy is 0.5565655048076923
Epoch 1500 Loss 0.6866980195045471
Training accuracy is 0.5565655048076923
Epoch 1550 Loss 0.6866980195045471
Training accuracy is 0.5565655048076923
Epoch 1600 Loss 0.6866980195045471
Training accuracy is 0.5565655048076923
Epoch 1650 Loss 0.686698317527771
Training accuracy is 0.5565655048076923
Epoch 1700 Loss 0.6866986155509949
Training accuracy is 0.5565655048076923
Epoch 1750 Loss 0.686698853969574
Training accuracy is 0.5565655048076923
Epoch 1800 Loss 0.6866990327835083
Training accuracy is 0.5565655048076923
Epoch 1850 Loss 0.6866992712020874
Training accuracy is 0.5565655048076923
Epoch 1900 Loss 0.6866998076438904
Training accuracy is 0.5565655048076923
Epoch 1950 Loss 0.6867008805274963
Training accuracy is 0.5565655048076923
Epoch 2000 Loss 0.6867018342018127
Training accuracy is 0.5565655048076923
Epoch 2050 Loss 0.6867029666900635
Training accuracy is 0.5565655048076923
Epoch 2100 Loss 0.6867037415504456
Training accuracy is 0.5565655048076923
Epoch 2150 Loss 0.68670654296875
Training accuracy is 0.5565655048076923
Epoch 2200 Loss 0.6867085099220276
Training accuracy is 0.5565655048076923
Epoch 2250 Loss 0.6867102980613708
Training accuracy is 0.5565655048076923
Epoch 2300 Loss 0.6867122650146484
Training accuracy is 0.5565655048076923
Epoch 2350 Loss 0.686714231967926
Training accuracy is 0.5565655048076923
Epoch 2400 Loss 0.6867154240608215
Training accuracy is 0.5565655048076923
Epoch 2450 Loss 0.6867167353630066
Training accuracy is 0.5565655048076923
Epoch 2500 Loss 0.6867179870605469
Training accuracy is 0.5565655048076923
Epoch 2550 Loss 0.6867192983627319
Training accuracy is 0.5565655048076923
Epoch 2600 Loss 0.6867207884788513
Training accuracy is 0.5565655048076923
Epoch 2650 Loss 0.6867217421531677
Training accuracy is 0.5565655048076923
Epoch 2700 Loss 0.6867228746414185
Training accuracy is 0.5565655048076923
Epoch 2750 Loss 0.6867242455482483
Training accuracy is 0.5565655048076923
Epoch 2800 Loss 0.6867255568504333
Training accuracy is 0.5565655048076923

Epoch 2850 Loss 0.6867268085479736
Training accuracy is 0.5565655048076923
Epoch 2900 Loss 0.6867278814315796
Training accuracy is 0.5565655048076923
Epoch 2950 Loss 0.6867291331291199
Training accuracy is 0.5565655048076923
Epoch 3000 Loss 0.6867303252220154
Training accuracy is 0.5565655048076923
Epoch 3050 Loss 0.6867312788963318
Training accuracy is 0.5565655048076923
Epoch 3100 Loss 0.6867325901985168
Training accuracy is 0.5565655048076923
Epoch 3150 Loss 0.6867333650588989
Training accuracy is 0.5565655048076923
Epoch 3200 Loss 0.6867353916168213
Training accuracy is 0.5565655048076923
Epoch 3250 Loss 0.6867369413375854
Training accuracy is 0.5565655048076923
Epoch 3300 Loss 0.6867368221282959
Training accuracy is 0.5565655048076923
Epoch 3350 Loss 0.6867377161979675
Training accuracy is 0.5565655048076923
Epoch 3400 Loss 0.6867396235466003
Training accuracy is 0.5565655048076923
Epoch 3450 Loss 0.6867396831512451
Training accuracy is 0.5565655048076923
Epoch 3500 Loss 0.6867386102676392
Training accuracy is 0.5565655048076923
Epoch 3550 Loss 0.686737596988678
Training accuracy is 0.5565655048076923
Epoch 3600 Loss 0.6867360472679138
Training accuracy is 0.5565655048076923
Epoch 3650 Loss 0.6867346167564392
Training accuracy is 0.5565655048076923
Epoch 3700 Loss 0.6867259740829468
Training accuracy is 0.5565655048076923
Epoch 3750 Loss 0.6867111921310425
Training accuracy is 0.5565655048076923
Epoch 3800 Loss 0.6866973042488098
Training accuracy is 0.5565655048076923
Epoch 3850 Loss 0.6866486072540283
Training accuracy is 0.5565655048076923
Epoch 3900 Loss 0.6865653395652771
Training accuracy is 0.5565655048076923
Epoch 3950 Loss 0.6864850521087646
Training accuracy is 0.5565655048076923
Epoch 4000 Loss 0.6862711310386658
Training accuracy is 0.5565655048076923
Epoch 4050 Loss 0.6855273842811584
Training accuracy is 0.5565655048076923
Epoch 4100 Loss 0.6563113331794739
Training accuracy is 0.5565655048076923
Epoch 4150 Loss 0.336140900850296
Training accuracy is 0.8275240384615384
Epoch 4200 Loss 0.19704784452915192
Training accuracy is 0.8982872596153846
Epoch 4250 Loss 0.3086509704589844

Training accuracy is 0.8722956730769231
Epoch 4300 Loss 0.2222747653722763
Training accuracy is 0.8900240384615384
Epoch 4350 Loss 0.2006864696741104
Training accuracy is 0.8977614182692307
Epoch 4400 Loss 0.2146114706993103
Training accuracy is 0.8958834134615384
Epoch 4450 Loss 0.1951487958431244
Training accuracy is 0.8981370192307693
Epoch 4500 Loss 0.19641393423080444
Training accuracy is 0.8979867788461539
Epoch 4550 Loss 0.2204727977514267
Training accuracy is 0.8942307692307693
Epoch 4600 Loss 0.1958731710910797
Training accuracy is 0.8974609375
Epoch 4650 Loss 0.20922620594501495
Training accuracy is 0.8915264423076923
Epoch 4700 Loss 0.21865598857402802
Training accuracy is 0.8641826923076923
Epoch 4750 Loss 0.23160573840141296
Training accuracy is 0.8965594951923077
Epoch 4800 Loss 0.1952981799840927
Training accuracy is 0.8982121394230769
Epoch 4850 Loss 0.20291200280189514
Training accuracy is 0.8958834134615384
Epoch 4900 Loss 0.19397442042827606
Training accuracy is 0.8997145432692307
Epoch 4950 Loss 0.19692093133926392
Training accuracy is 0.8975360576923077
Epoch 5000 Loss 0.20486140251159668
Training accuracy is 0.8934795673076923

```
In [26]: # Further training is needed
# Also, the learning rate will be changed to 0.05
# Based on the outputs above, stop training when accuracy reaches 0.899
learning_rate = 0.05
optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate, momentum=0.9)
for epoch in range(nepochs):
    optimizer.zero_grad()
    net_out = net(data)
    loss = criterion(net_out, target)
    loss.backward()
    optimizer.step()

    pred = net_out.data.max(1)[1]
    correctidx = pred.eq(target.data)
    ncorrect = correctidx.sum()
    accuracy = ncorrect.item()/len(data)
    if accuracy >= 0.899:
        print('Epoch ', epoch, 'Loss ', loss.item())
        print('Training accuracy is ', accuracy)
        break
```

C:\Anaconda3\lib\site-packages\ipykernel_launcher.py:34: UserWarning: Implicit dimension choice for log_softmax has been deprecated. Change the call to include dim=X as an argument.

```
Epoch 45 Loss 0.19511112570762634
Training accuracy is 0.8993389423076923
```

Report:

```
In [28]: print('Accuracy:', accuracy, '\n' + 'Loss:', loss.item())
print('Number of hidden layers: 10; type of activation function: relu; number
of neurons per layer: 20')


# plot the decision boundary of the network
plot_decision_boundary(net, X, y)
```

Accuracy: 0.8993389423076923

Loss: 0.19511112570762634

Number of hidden layers: 10; type of activation function: relu; number of neurons per layer: 20

C:\Anaconda3\lib\site-packages\ipykernel_launcher.py:34: UserWarning: Implicit dimension choice for log_softmax has been deprecated. Change the call to include dim=X as an argument.




2.2

Important note: This dataset is class-imbalanced, so I am using the oversampling method to prevent the classification result from being dominated by the majority class.

```
In [3]: mydata2 = pd.read_csv('FeedForward_Data_hexa.csv', header=None)
X = mydata2.values[:, 0:2] # features
y = mydata2.values[:, 2] # labels

# plot the original data
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Spectral, s = 1)
```

Out[3]: <matplotlib.collections.PathCollection at 0x1d1043a6e48>



```
In [4]: # Oversampling the rare cases (y=0) with replacement
# Note that the number of majority cases = 12521 while the number of rare case
s = 791
mydata2.columns = ['x1', 'x2', 'y']
sample = mydata2[mydata2['y'] == 0].sample(n=11730, replace=True)
# add the new samples to X and y
new_X = np.concatenate((X, sample.values[:, 0:2]))
new_y = np.concatenate((y, sample.values[:, 2]))
new_X = torch.tensor(new_X, dtype = torch.float)
new_y = torch.tensor(new_y, dtype = torch.long)
```

```

In [8]: class Net(nn.Module):
        # network structure for 2.2
        def __init__(self):
            super(Net, self).__init__()
            # 10 hidden layers with 30 neurons each
            self.fc1 = nn.Linear(2, 30)
            self.fc2 = nn.Linear(30, 30)
            self.fc3 = nn.Linear(30, 30)
            self.fc4 = nn.Linear(30, 30)
            self.fc5 = nn.Linear(30, 30)
            self.fc6 = nn.Linear(30, 30)
            self.fc7 = nn.Linear(30, 30)
            self.fc8 = nn.Linear(30, 30)
            self.fc9 = nn.Linear(30, 30)
            self.fc10 = nn.Linear(30, 30)
            self.fc11 = nn.Linear(30, 30)
            self.fc12 = nn.Linear(30, 2)

        def forward(self, x):
            # choice of activation function: relu, tanh
            x = torch.relu(self.fc1(x))
            x = torch.relu(self.fc2(x))
            x = torch.relu(self.fc3(x))
            x = torch.relu(self.fc4(x))
            x = torch.relu(self.fc5(x))
            x = torch.relu(self.fc6(x))
            x = torch.tanh(self.fc7(x))
            x = torch.tanh(self.fc8(x))
            x = torch.tanh(self.fc9(x))
            x = torch.tanh(self.fc10(x))
            x = torch.tanh(self.fc11(x))
            x = self.fc12(x)
            return F.log_softmax(x)

# Training
net = Net()
learning_rate = 0.05
optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate, momentum=0.9)
criterion = nn.NLLLoss()
nepochs = 10001
data, target = new_X, new_y
for epoch in range(nepochs):
    optimizer.zero_grad()
    # forward propagate
    net_out = net(data)
    # compute loss
    loss = criterion(net_out, target)
    # backpropagate
    loss.backward()
    # update parameters
    optimizer.step()

    pred = net_out.data.max(1)[1]
    correctidx = pred.eq(target.data)
    ncorrect = correctidx.sum()
    accuracy = ncorrect.item()/len(data)

```



```
if accuracy >= 0.95:
    print('Epoch ', epoch, 'Loss ', loss.item())
    print('Training accuracy is ', accuracy)
    break
if epoch % 100 == 0:
    print('Epoch ', epoch, 'Loss ', loss.item())
    print('Training accuracy is ', accuracy)
```

```
C:\Anaconda3\lib\site-packages\ipykernel_launcher.py:33: UserWarning: Implicit dimension choice for log_softmax has been deprecated. Change the call to include dim=X as an argument.
```

Epoch 0 Loss 0.6978422999382019
Training accuracy is 0.5
Epoch 100 Loss 0.6930481195449829
Training accuracy is 0.5
Epoch 200 Loss 0.693304181098938
Training accuracy is 0.6009504033224183
Epoch 300 Loss 0.6933044791221619
Training accuracy is 0.5792668317227059
Epoch 400 Loss 0.6933032870292664
Training accuracy is 0.592364827090488
Epoch 500 Loss 0.6933022141456604
Training accuracy is 0.5911668397092884
Epoch 600 Loss 0.6933016777038574
Training accuracy is 0.5909671751457551
Epoch 700 Loss 0.6932948231697083
Training accuracy is 0.6063812794505231
Epoch 800 Loss 0.6932891011238098
Training accuracy is 0.6241913585176903
Epoch 900 Loss 0.6932847499847412
Training accuracy is 0.6409631818544844
Epoch 1000 Loss 0.6932802200317383
Training accuracy is 0.6423209008865106
Epoch 1100 Loss 0.6932776570320129
Training accuracy is 0.6558980912067727
Epoch 1200 Loss 0.6932759881019592
Training accuracy is 0.6766632058142321
Epoch 1300 Loss 0.6932615637779236
Training accuracy is 0.6968692596437984
Epoch 1400 Loss 0.6932504177093506
Training accuracy is 0.7015414104304768
Epoch 1500 Loss 0.6930919885635376
Training accuracy is 0.7129622234645795
Epoch 1600 Loss 0.6928747892379761
Training accuracy is 0.739357878763677
Epoch 1700 Loss 0.6926199197769165
Training accuracy is 0.7467454676144077
Epoch 1800 Loss 0.6909545660018921
Training accuracy is 0.7634773580384954
Epoch 1900 Loss 0.7013331055641174
Training accuracy is 0.5
Epoch 2000 Loss 0.6931933164596558
Training accuracy is 0.5
Epoch 2100 Loss 0.6932984590530396
Training accuracy is 0.8106381279450523
Epoch 2200 Loss 0.6932749152183533
Training accuracy is 0.7208290072677901
Epoch 2300 Loss 0.693254292011261
Training accuracy is 0.6902004632217874
Epoch 2400 Loss 0.6932433247566223
Training accuracy is 0.7002635572238639
Epoch 2500 Loss 0.6932215094566345
Training accuracy is 0.7522562095679259
Epoch 2600 Loss 0.6931699514389038
Training accuracy is 0.7896733487740596
Epoch 2700 Loss 0.6928461790084839
Training accuracy is 0.8253733727338072
Epoch 2800 Loss 0.6914950609207153

Training accuracy is 0.8367941857679099
Epoch 2900 Loss 0.7084758281707764
Training accuracy is 0.5
Epoch 3000 Loss 0.7805532217025757
Training accuracy is 0.5
Epoch 3100 Loss 0.693059504032135
Training accuracy is 0.5
Epoch 3200 Loss 0.6932235956192017
Training accuracy is 0.694273620317866
Epoch 3300 Loss 0.6929205060005188
Training accuracy is 0.7138806804568325
Epoch 3400 Loss 0.6925575733184814
Training accuracy is 0.7237840428080824
Epoch 3500 Loss 0.6717435121536255
Training accuracy is 0.7238639086334957
Epoch 3600 Loss 0.4430851340293884
Training accuracy is 0.7966216755850172
Epoch 3700 Loss 0.5165365934371948
Training accuracy is 0.8292069323536458
Epoch 3800 Loss 0.2585148811340332
Training accuracy is 0.9193754492452679
Epoch 3811 Loss 0.16301041841506958
Training accuracy is 0.95204057183931

Report:

```
In [9]: print('Accuracy:', accuracy, 'Loss:', loss.item())
print('I used 10 hidden layers with 30 neurons per layer.')
print('Type of activation functions: relu and tanh.')


# plot the decision boundary of the network (using the original X and y)
plot_decision_boundary(net, X, y)
```

Accuracy: 0.95204057183931 Loss: 0.16301041841506958

I used 10 hidden layers with 30 neurons per layer.

Type of activation functions: relu and tanh.

C:\Anaconda3\lib\site-packages\ipykernel_launcher.py:33: UserWarning: Implicit dimension choice for log_softmax has been deprecated. Change the call to include dim=X as an argument.



At first, my network could achieve 94% accuracy very quickly, and that was 'fake' since the network did nothing but classified all the points as '1'. Then I applied oversampling method and an accuracy value significantly higher than 50% should be acceptable because it is actually separating the data points. Based on the plot above, I think it did a good job.